

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature		Description
<code>project_id</code>		A unique identifier for the proposed project. Example: p036502
<code>project_title</code>	<ul style="list-style-type: none">••	Title of the project. Examples: Art Will Make You Happy! First Grade Fun
<code>project_grade_category</code>	<ul style="list-style-type: none">••••	Grade level of students for which the project is targeted. One of the following enumerated values: Grades PreK-2 Grades 3-5 Grades 6-8 Grades 9-12
<code>project_subject_categories</code>	<ul style="list-style-type: none">••••••••	One or more (comma-separated) subject categories for the project from the following enumerated list of values: Applied Learning Care & Hunger Health & Sports History & Civics Literacy & Language Math & Science Music & The Arts Special Needs Warmth
	<ul style="list-style-type: none">••	Examples: Music & The Arts Literacy & Language, Math & Science
<code>school_state</code>		State where school is located (Two-letter U.S. postal code (https://en.wikipedia.org/wiki/List_of_U.S._state_abbreviations#Postal_codes)). Example: WY
<code>project_subject_subcategories</code>	<ul style="list-style-type: none">••	One or more (comma-separated) subject subcategories for the project. Examples: Literacy Literature & Writing, Social Sciences
<code>project_resource_summary</code>	<ul style="list-style-type: none">•	An explanation of the resources needed for the project. Example: My students need hands on literacy materials to manage sensory needs!</code>
<code>project_essay_1</code>		First application essay*
<code>project_essay_2</code>		Second application essay*
<code>project_essay_3</code>		Third application essay*
<code>project_essay_4</code>		Fourth application essay*
<code>project_submitted_datetime</code>		Datetime when project application was submitted. Example: 2016-04-28 12:43:56.245
<code>teacher_id</code>		A unique identifier for the teacher of the proposed project. Example: bdf8baa8fedef6bfeec7ae4ff1c15c56
<code>teacher_prefix</code>	<ul style="list-style-type: none">••••••	Teacher's title. One of the following enumerated values: nan Dr. Mr. Mrs. Ms. Teacher.
<code>teacher_number_of_previously_posted_projects</code>		Number of project applications previously submitted by the same teacher. Example: 2

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
<code>id</code>	A <code>project_id</code> value from the <code>train.csv</code> file. Example: <code>p036502</code>
<code>description</code>	Description of the resource. Example: <code>Tenor Saxophone Reeds, Box of 25</code>
<code>quantity</code>	Quantity of the resource required. Example: <code>3</code>
<code>price</code>	Price of the resource required. Example: <code>9.95</code>

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
<code>project_is_approved</code>	A binary flag indicating whether DonorsChoose approved the project. A value of <code>0</code> indicates the project was not approved, and a value of <code>1</code> indicates the project was approved.

Notes on the Essay Data

- Prior to May 17, 2016, the prompts for the essays were as follows:
- **project_essay_1:** "Introduce us to your classroom"
 - **project_essay_2:** "Tell us more about your students"
 - **project_essay_3:** "Describe how your students will use the materials you're requesting"
 - **project_essay_3:** "Close by sharing why your project will make a difference"
- Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:
- **project_essay_1:** "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
 - **project_essay_2:** "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be `NaN`.

```

In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
import math
from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_validate

from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import LogisticRegression

```

1.1 Reading Data

```

In [30]: project_data = pd.read_csv('train_data.csv')
resource_data = pd.read_csv('resources.csv')

project_data=project_data.dropna(subset=['teacher_prefix'])

```

```

In [31]: print("Number of data points in train data", project_data.shape)
print('-'*50)
print("The attributes of data :", project_data.columns.values)

```

Number of data points in train data (109245, 17)

The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state' 'project_submitted_datetime' 'project_grade_category' 'project_subject_categories' 'project_subject_subcategories' 'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3' 'project_essay_4' 'project_resource_summary' 'teacher_number_of_previously_posted_projects' 'project_is_approved']

```

In [32]: print("Number of data points in train data", resource_data.shape)
print(resource_data.columns.values)
resource_data.head(2)

```

Number of data points in train data (1541272, 4)

['id' 'description' 'quantity' 'price']

Out[32]:

	id	description	quantity	price
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	149.00
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	14.95

1.2 preprocessing of project_subject_categories

```
In [34]: categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math","&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
        j = j.replace(' ','') # we are replacing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&Science"
        temp+=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&','_') # we are replacing the & value into _
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))
```

1.3 preprocessing of project_subject_subcategories

```
In [35]: sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math","&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
        j = j.replace(' ','') # we are replacing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&Science"
        temp +=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))
```

1.3 Text preprocessing

```
In [36]: # merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)
```

```
project_data.head(2)
```

Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_submitted_datetime	project_grade_category	project_title	
0	160221	p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs.	IN	2016-12-05 13:43:57	Grades PreK-2	Educational Support English Learner Home
1	140945	p258326	897464ce9ddc600bcd1151f324dd63a	Mr.	FL	2016-10-25 09:22:10	Grades 6-8	Wanted Project Hurricane Learning

```
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"can't", "can not", phrase, flags= re.IGNORECASE)

    # general
    phrase = re.sub(r"n't", " not", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'re", " are", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'s", " is", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'d", " would", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'ll", " will", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'t", " not", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'ve", " have", phrase, flags= re.IGNORECASE)
    phrase = re.sub(r"\'m", " am", phrase, flags= re.IGNORECASE)
    return phrase
```

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"]
```


```
# Combining all the above students
from tqdm import tqdm
preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    preprocessed_essays.append(sent.lower().strip())
```

```
100%|██████████████████████████████████████████████████████████████████████████████| 109245/109245 [01:05<00:00, 1663.76it/s]
```

```
# after preprocessing
preprocessed_essays[20000]
```

'my kindergarten students varied disabilities ranging speech language delays cognitive delays gross fine motor delays autism they eager beavers always strive work hardest working past limitations the materials ones i seek students i teach title i school students receive free reduced price lunch despite disabilities limitations students love coming school come eager learn explore have ever felt like ants pants needed groove move meeting this kids feel time the want able move learn say wobble chairs answer i love develop core enhances gross motor turn fine motor skills they also want learn games kids not want sit worksheets they want learn count jumping playing physical engagement key success the number toss color shape mats make happen my students forget work fun 6 year old deserves nannan'

- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



- Along with plotting ROC curve, you need to print the [confusion matrix](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>), with predicted and original labels of test data points. Please visualize your confusion matrices using [seaborn heatmaps](#).

	Predicted: NO	Predicted: YES
Actual: NO	TN = ??	FP = ??
Actual: YES	FN = ??	TP = ??

(<https://seaborn.pydata.org/generated/seaborn.heatmap.html>).

4. [Task-2] Apply Logistic Regression on the below feature set **Set 5** by finding the best hyper parameter as suggested in step 2 and step 3.
5. Consider these set of features **Set 5** :

- school_state** : categorical data
- clean_categories** : categorical data
- clean_subcategories** : categorical data
- project_grade_category** :categorical data
- teacher_prefix** : categorical data
- quantity** : numerical data
- teacher_number_of_previously_posted_projects** : numerical data
- price** : numerical data
- sentiment score's of each of the essay** : numerical data
- number of words in the title** : numerical data
- number of words in the combine essays** : numerical data

And apply the Logistic regression on these features by finding the best hyper paramter as suggested in step 2 and step 3

6. Conclusion

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library [link](http://zetcode.com/python/prettytable/) (<http://zetcode.com/python/prettytable/>).

Vectorizer	Model	Hyper parameter	AUC
BOW	Brute	7	0.78
TFIDF	Brute	12	0.79
W2V	Brute	10	0.78
TFIDFW2V	Brute	6	0.78

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this [link](https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf). (<https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf>)

2. Logistic Regression

2.1 Splitting data into Train and cross validation(or test): Stratified Sampling

```
In [44]: data= project_data

y = data['project_is_approved'].values
data.drop(['project_is_approved'], axis=1, inplace=True)
x=data

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, stratify=y)

print(x_train.shape, y_train.shape)
print(x_test.shape, y_test.shape)

(76471, 22) (76471,)
(32774, 22) (32774,)
```

2.2 Make Data Model Ready: encoding numerical, categorical features

Vectorizing Data and Creating Data Matrix

Vectorizing categorical features:

```
In [45]: # Vectorizing categorical features:

#school_state

vectorizer_school_state = CountVectorizer()

# we use the fitted CountVectorizer to convert the text to vector
x_train_state_ohe = vectorizer_school_state.fit_transform(x_train['school_state'].values) # fit has to happen only on tra
x_test_state_ohe = vectorizer_school_state.transform(x_test['school_state'].values)

print("After vectorizations")
print(x_train_state_ohe.shape, y_train.shape)
print(x_test_state_ohe.shape, y_test.shape)
print(vectorizer_school_state.get_feature_names())
print("="*100)

After vectorizations
(76471, 51) (76471,)
(32774, 51) (32774,)
['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl', 'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'm
a', 'md', 'me', 'mi', 'mn', 'mo', 'ms', 'mt', 'nc', 'nd', 'ne', 'nh', 'nj', 'nm', 'nv', 'ny', 'oh', 'ok', 'or', 'pa',
'ri', 'sc', 'sd', 'tn', 'tx', 'ut', 'va', 'vt', 'wa', 'wi', 'wv', 'wy']
=====
```

```
In [46]: # Vectorizing categorical features:

#clean_categories

vectorizer_cat = CountVectorizer()

# we use the fitted CountVectorizer to convert the text to vector
x_train_clean_cat_ohe = vectorizer_cat.fit_transform(x_train['clean_categories'].values) # fit has to happen only on trai
x_test_clean_cat_ohe = vectorizer_cat.transform(x_test['clean_categories'].values)

print("After vectorizations")
print(x_train_clean_cat_ohe.shape, y_train.shape)
print(x_test_clean_cat_ohe.shape, y_test.shape)
print(vectorizer_cat.get_feature_names())
print("="*100)

After vectorizations
(76471, 9) (76471,)
(32774, 9) (32774,)
['appliedlearning', 'care_hunger', 'health_sports', 'history_civics', 'literacy_language', 'math_science', 'music_art
s', 'specialneeds', 'warmth']
=====
```



```
In [47]: # Vectorizing categorical features:

#clean_subcategories

vectorizer_subcat = CountVectorizer()

# we use the fitted CountVectorizer to convert the text to vector
x_train_clean_subcat_ohe = vectorizer_subcat.fit_transform(x_train['clean_subcategories'].values) # fit has to happen onl
x_test_clean_subcat_ohe = vectorizer_subcat.transform(x_test['clean_subcategories'].values)

print("After vectorizations")
print(x_train_clean_subcat_ohe.shape, y_train.shape)
print(x_test_clean_subcat_ohe.shape, y_test.shape)
print(vectorizer_subcat.get_feature_names())
print("="*100)

After vectorizations
(76471, 30) (76471,)
(32774, 30) (32774,)
['appliedsciences', 'care_hunger', 'charactereducation', 'civics_government', 'college_careerprep', 'communityservice',
'earlydevelopment', 'economics', 'environmentalscience', 'esl', 'extracurricular', 'financialliteracy', 'foreignlanguag
es', 'gym_fitness', 'health_lifescience', 'health_wellness', 'history_geography', 'literacy', 'literature_writing', 'ma
thematics', 'music', 'nutritioneducation', 'other', 'parentinvolvement', 'performingarts', 'socialsciences', 'specialne
eds', 'teamsports', 'visualarts', 'warmth']
=====
```

```
In [48]: # Vectorizing categorical features:

#project_grade_category

vectorizer_grade = CountVectorizer()

# we use the fitted CountVectorizer to convert the text to vector
x_train_grade_ohe = vectorizer_grade.fit_transform(x_train['project_grade_cat_cleaned'].values) # fit has to happen only
x_test_grade_ohe = vectorizer_grade.transform(x_test['project_grade_cat_cleaned'].values)

print("After vectorizations")
print(x_train_grade_ohe.shape, y_train.shape)
print(x_test_grade_ohe.shape, y_test.shape)
print(vectorizer_grade.get_feature_names())
print("="*100)

After vectorizations
(76471, 4) (76471,)
(32774, 4) (32774,)
['grades_3_to_5', 'grades_6_to_8', 'grades_9_to_12', 'grades_prek_to_2']
=====
```

```
In [49]: # Vectorizing categorical features:

#teacher_prefix

vectorizer_prefix = CountVectorizer()

# we use the fitted CountVectorizer to convert the text to vector
x_train_teacher_ohe = vectorizer_prefix.fit_transform(x_train['teacher_prefix'].values) # fit has to happen only on train
x_test_teacher_ohe = vectorizer_prefix.transform(x_test['teacher_prefix'].values)

print("After vectorizations")
print(x_train_teacher_ohe.shape, y_train.shape)
print(x_test_teacher_ohe.shape, y_test.shape)
print(vectorizer_prefix.get_feature_names())
print("="*100)

After vectorizations
(76471, 5) (76471,)
(32774, 5) (32774,)
['dr', 'mr', 'mrs', 'ms', 'teacher']
=====
```

Normalizing numerical features:

In [50]: # Normalizing numerical features:

#price

```
from sklearn.preprocessing import Normalizer
normalizer_price = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
# normalizer.fit(X_train['price'].values.reshape(-1,1))

x_train_price_norm = normalizer_price.fit_transform(x_train['price'].values.reshape(-1,1))
x_test_price_norm = normalizer_price.transform(x_test['price'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_price_norm.shape, y_train.shape)
print(x_test_price_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(76471, 1) (76471,)

(32774, 1) (32774,)

=====

In [51]: # Normalizing numerical features:

#teacher_number_of_previously_posted_projects

```
from sklearn.preprocessing import Normalizer
normalizer_post_proj = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
# normalizer.fit(X_train['price'].values.reshape(-1,1))

x_train_post_proj_norm = normalizer_post_proj.fit_transform(x_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
x_test_post_proj_norm = normalizer_post_proj.transform(x_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_post_proj_norm.shape, y_train.shape)
print(x_test_post_proj_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(76471, 1) (76471,)

(32774, 1) (32774,)

=====

In [52]: # Normalizing numerical features:

#quantity

```
from sklearn.preprocessing import Normalizer
normalizer_quantity = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
# normalizer.fit(X_train['price'].values.reshape(-1,1))

x_train_quantity_norm = normalizer_quantity.fit_transform(x_train['quantity'].values.reshape(-1,1))
x_test_quantity_norm = normalizer_quantity.transform(x_test['quantity'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_quantity_norm.shape, y_train.shape)
print(x_test_quantity_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(76471, 1) (76471,)

(32774, 1) (32774,)

=====

2.3 Make Data Model Ready: encoding eassay, and project_title

Bag of Words(BoW) : Essay

```
In [53]: vectorizer_essay_bow = CountVectorizer(min_df=10,ngram_range=(2,2), max_features=5000)

x_train_essay_bow = vectorizer_essay_bow.fit_transform(x_train['preprocessed_essays'].values)
x_test_essay_bow = vectorizer_essay_bow.transform(x_test['preprocessed_essays'].values)

print("After vectorizations")
print(x_train_essay_bow.shape, y_train.shape)
print(x_test_essay_bow.shape, y_test.shape)
print("="*100)

After vectorizations
(76471, 5000) (76471,)
(32774, 5000) (32774,)
=====
```

Bag of Words(BoW) : Project_title

```
In [54]: vectorizer_title_bow = CountVectorizer(min_df=0)

x_train_title_bow = vectorizer_title_bow.fit_transform(x_train['preprocessed_project_title'].values)
x_test_title_bow = vectorizer_title_bow.transform(x_test['preprocessed_project_title'].values)

print("After vectorizations")
print(x_train_title_bow.shape, y_train.shape)
print(x_test_title_bow.shape, y_test.shape)
print("="*100)

After vectorizations
(76471, 14392) (76471,)
(32774, 14392) (32774,)
=====
```

TFIDF Vectorizer: Essay

```
In [55]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer_essay_tfidf = TfidfVectorizer(min_df=10,ngram_range=(2,2),max_features=5000)

x_train_essay_tfidf = vectorizer_essay_tfidf.fit_transform(x_train['preprocessed_essays'].values)
x_test_essay_tfidf = vectorizer_essay_tfidf.transform(x_test['preprocessed_essays'].values)

print("After vectorizations")
print(x_train_essay_tfidf.shape, y_train.shape)
print(x_test_essay_tfidf.shape, y_test.shape)
print("="*100)

After vectorizations
(76471, 5000) (76471,)
(32774, 5000) (32774,)
=====
```

TFIDF Vectorizer: Project_title

```
In [56]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer_title_tfidf = TfidfVectorizer(min_df=0)

x_train_title_tfidf = vectorizer_title_tfidf.fit_transform(x_train['preprocessed_project_title'].values)
x_test_title_tfidf = vectorizer_title_tfidf.transform(x_test['preprocessed_project_title'].values)

print("After vectorizations")
print(x_train_title_tfidf.shape, y_train.shape)
print(x_test_title_tfidf.shape, y_test.shape)
print("="*100)

After vectorizations
(76471, 14392) (76471,)
(32774, 14392) (32774,)
=====
```

Avg W2V Vectorizer: Essay

```
In [57]: # stronging variables into pickle files python: http://www.jessicayung.com/how-to-use-pickle-to-save-and-load-variables-i
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

```
In [58]: # average Word2Vec
# compute average word2vec for each review.
train_w2v_vectors_essays = []; # the avg-w2v for each essay is stored in this list
for sentence in tqdm(x_train['preprocessed_essays'].values): # for each essay in training data
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the essay
    for word in sentence.split(): # for each word in a essay
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    train_w2v_vectors_essays.append(vector)
print("train vector")
print(len(train_w2v_vectors_essays))
print(len(train_w2v_vectors_essays[0]))
print('='*50)

test_w2v_vectors_essays = []; # the avg-w2v for each essay is stored in this list
for sentence in tqdm(x_test['preprocessed_essays'].values): # for each essay in training data
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the essay
    for word in sentence.split(): # for each word in a essay
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    test_w2v_vectors_essays.append(vector)
print("test vector")
print(len(test_w2v_vectors_essays))
print(len(test_w2v_vectors_essays[0]))
print('='*50)
```

[illegible]

```
train  vector
76471
300
```

[illegible]

```
test  vector
32774
300
```

Avg W2V Vectorizer: Project_Title

TFIDF W2V Vectorizer: Essay

13/28


```
In [68]: #AVG W2V

from scipy.sparse import hstack
x_w2v_tr = hstack((x_train_state_oh, x_train_clean_cat_oh, x_train_clean_subcat_oh, x_train_grade_oh, x_train_teacher_oh,
                  x_train_price_norm, x_train_post_proj_norm, x_train_quantity_norm, \
                  train_w2v_vectors_essays, train_w2v_vectors_title)).tocsr()

x_w2v_te = hstack((x_test_state_oh, x_test_clean_cat_oh, x_test_clean_subcat_oh, x_test_grade_oh, x_test_teacher_oh,
                  x_test_price_norm, x_test_post_proj_norm, x_test_quantity_norm, \
                  test_w2v_vectors_essays, test_w2v_vectors_title)).tocsr()
```

```
In [69]: # TFDIF AVG W2V

from scipy.sparse import hstack
x_tfidf_w2v_tr = hstack((x_train_state_oh, x_train_clean_cat_oh, x_train_clean_subcat_oh, x_train_grade_oh, x_train_teacher_oh,
                        x_train_price_norm, x_train_post_proj_norm, x_train_quantity_norm, \
                        train_tfidf_w2v_essays, train_tfidf_w2v_title)).tocsr()

x_tfidf_w2v_te = hstack((x_test_state_oh, x_test_clean_cat_oh, x_test_clean_subcat_oh, x_test_grade_oh, x_test_teacher_oh,
                        x_test_price_norm, x_test_post_proj_norm, x_test_quantity_norm, \
                        test_tfidf_w2v_essays, test_tfidf_w2v_title)).tocsr()
```

2.4 Applying Logistic Regression on different kind of featurization as mentioned in the instructions

Apply Logistic Regression on different kind of featurization as mentioned in the instructions
For Every model that you work on make sure you do the step 2 and step 3 of instructions

2.4.1 Applying LR brute force on BOW, SET 1

```
In [63]: C=[10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]

tuned_parameters = {'C': C}

tuned_parameters2 = {'C': [math.log(i) for i in C]}


model_lr = RandomizedSearchCV(LogisticRegression(), tuned_parameters, cv=3, scoring='roc_auc')
model_lr.fit(x_bow_tr, y_train)

train_auc= model_lr.cv_results_['mean_train_score']
train_auc_std= model_lr.cv_results_['std_train_score']
cv_auc = model_lr.cv_results_['mean_test_score']
cv_auc_std= model_lr.cv_results_['std_test_score']

plt.plot(tuned_parameters2['C'], train_auc, label='Train AUC')
plt.plot(tuned_parameters2['C'], cv_auc, label='CV AUC')

plt.scatter(tuned_parameters2['C'], train_auc, label='Train AUC points')
plt.scatter(tuned_parameters2['C'], cv_auc, label='CV AUC points')

plt.grid()
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- BAG OF WORDS")
plt.show()
```



Now Train your Model using Hyper-parameter

```
In [134]: from sklearn.metrics import roc_curve, auc, roc_auc_score

model_lr = LogisticRegression(C=0.01,class_weight='balanced')
model_lr.fit(x_bow_tr, y_train) # Training Data


y_train_pred = LogisticRegression.predict_proba(model_lr,x_bow_tr)[:,-1] # Predicting Train Data
y_test_pred = LogisticRegression.predict_proba(model_lr,x_bow_te)[:,-1] # Predicting Test Data based on training model

y_train_pred_con = LogisticRegression.predict(model_lr,x_bow_tr)
y_test_pred_con = LogisticRegression.predict(model_lr,x_bow_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred) # GET TRAIN TRUE POSITIVE and FALSE POSITIVE RATES
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred) # GET TEST TRUE POSITIVE and FALSE POSITIVE RATES

plt.grid()
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- BoW")
plt.show()

auc_bow=auc(test_fpr, test_tpr)
```




```
In [123]: #function to get heatmap confusion matrix-- source kaggle https://www.kaggle.com/shashank49/donors-choose-knn
def get_confusion_matrix(y_train,y_train_pred):
    df_cm = pd.DataFrame(confusion_matrix(y_train,y_train_pred), range(2),range(2))
    df_cm.columns = ['Predicted NO', 'Predicted YES']
    df_cm = df_cm.rename({0: 'Actual NO', 1: 'Actual YES'})
    sns.set(font_scale=1.4)#for label size
    sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
In [136]: from sklearn.metrics import confusion_matrix

print("Train confusion matrix")
get_confusion_matrix(y_train,y_train_pred_con)
print('='*100)
```

Train confusion matrix


=====



```
In [137]: print("Test confusion matrix")
print('='*100)
get_confusion_matrix(y_test,y_test_pred_con)
```

Test confusion matrix

=====



2.4.2 Applying LR brute force on TFIDF, SET 2

```
In [65]: C=[10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]

tuned_parameters = {'C': C}

tuned_parameters2 = {'C': [math.log(i) for i in C]}


model_lr = RandomizedSearchCV(LogisticRegression(), tuned_parameters, cv=3, scoring='roc_auc')
model_lr.fit(x_tfidf_tr, y_train)

train_auc= model_lr.cv_results_['mean_train_score']
train_auc_std= model_lr.cv_results_['std_train_score']
cv_auc = model_lr.cv_results_['mean_test_score']
cv_auc_std= model_lr.cv_results_['std_test_score']

plt.plot(tuned_parameters2['C'], train_auc, label='Train AUC')
plt.plot(tuned_parameters2['C'], cv_auc, label='CV AUC')

plt.scatter(tuned_parameters2['C'], train_auc, label='Train AUC points')
plt.scatter(tuned_parameters2['C'], cv_auc, label='CV AUC points')

plt.grid()
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- TFIDF")
plt.show()
```



Now Train your Model using Hyper-parameter

```
In [138]: from sklearn.metrics import roc_curve, auc

model_lr = LogisticRegression(C=0.1,class_weight='balanced')
model_lr.fit(x_tfidf_tr, y_train) # Training Data


y_train_pred = LogisticRegression.predict_proba(model_lr, x_tfidf_tr)[:,-1] # Predicting Train Data
y_test_pred = LogisticRegression.predict_proba(model_lr, x_tfidf_te)[:,-1] # Predicting Test Data based on training mo

y_train_pred_con = LogisticRegression.predict(model_lr, x_tfidf_tr)
y_test_pred_con = LogisticRegression.predict(model_lr, x_tfidf_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred) # GET TRAIN TRUE POSITIVE and FALSE POSITIVE RATES
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred) # GET TEST TRUE POSITIVE and FALSE POSITIVE RATES

plt.grid()
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- TFIDF")
plt.show()

auc_tfidf=auc(test_fpr, test_tpr)
```




```
In [139]: from sklearn.metrics import confusion_matrix

print("Train confusion matrix")
get_confusion_matrix(y_train,y_train_pred_con)
print('='*100)
```

Train confusion matrix


=====



```
In [140]: print("Test confusion matrix")
print('='*100)
get_confusion_matrix(y_test,y_test_pred_con)
```

Test confusion matrix

=====



2.4.3 Applying LR brute force on W2V, SET 3

```
In [66]: C=[10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]

tuned_parameters = {'C': C}

tuned_parameters2 = {'C': [math.log(i) for i in C]}


model_lr = RandomizedSearchCV(LogisticRegression(), tuned_parameters, cv=3, scoring='roc_auc')
model_lr.fit(x_w2v_tr, y_train)

train_auc= model_lr.cv_results_['mean_train_score']
train_auc_std= model_lr.cv_results_['std_train_score']
cv_auc = model_lr.cv_results_['mean_test_score']
cv_auc_std= model_lr.cv_results_['std_test_score']

plt.plot(tuned_parameters2['C'], train_auc, label='Train AUC')
plt.plot(tuned_parameters2['C'], cv_auc, label='CV AUC')

plt.scatter(tuned_parameters2['C'], train_auc, label='Train AUC points')
plt.scatter(tuned_parameters2['C'], cv_auc, label='CV AUC points')

plt.grid()
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- W2V")
plt.show()
```



Now Train your Model using Hyper-parameter


```
In [141]: from sklearn.metrics import roc_curve, auc

model_lr = LogisticRegression(C=1,class_weight='balanced')
model_lr.fit(x_w2v_tr, y_train) # Training Data


y_train_pred = LogisticRegression.predict_proba(model_lr, x_w2v_tr)[:,-1] # Predicting Train Data
y_test_pred = LogisticRegression.predict_proba(model_lr, x_w2v_te) [-,-1] # Predicting Test Data based on training mode

y_train_pred_con = LogisticRegression.predict(model_lr, x_w2v_tr) # Predicting Train Data
y_test_pred_con = LogisticRegression.predict(model_lr, x_w2v_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred) # GET TRAIN TRUE POSITIVE and FALSE POSITIVE RATES
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred) # GET TEST TRUE POSITIVE and FALSE POSITIVE RATES


plt.grid()
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- W2V")
plt.show()

auc_w2v=auc(test_fpr, test_tpr)
```



```
In [142]: from sklearn.metrics import confusion_matrix


print("Train confusion matrix")
get_confusion_matrix(y_train,y_train_pred_con)
print('='*100)
```



```
In [143]: print("Test confusion matrix")
print('='*100)
get_confusion_matrix(y_test,y_test_pred_con)
```

Test confusion matrix

=====



2.4.4 Applying LR brute force on TFIDF W2V, SET 4

```
In [67]: C=[10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]

tuned_parameters = {'C': C}

tuned_parameters2 = {'C': [math.log(i) for i in C]}


model_lr = RandomizedSearchCV(LogisticRegression(), tuned_parameters, cv=3, scoring='roc_auc')
model_lr.fit(x_tfidf_w2v_tr, y_train)

train_auc= model_lr.cv_results_['mean_train_score']
train_auc_std= model_lr.cv_results_['std_train_score']
cv_auc = model_lr.cv_results_['mean_test_score']
cv_auc_std= model_lr.cv_results_['std_test_score']

plt.plot(tuned_parameters2['C'], train_auc, label='Train AUC')
plt.plot(tuned_parameters2['C'], cv_auc, label='CV AUC')

plt.scatter(tuned_parameters2['C'], train_auc, label='Train AUC points')
plt.scatter(tuned_parameters2['C'], cv_auc, label='CV AUC points')

plt.grid()
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- TFIDF W2v")
plt.show()
```



Now Train your Model using Hyper-parameter

```
In [144]: from sklearn.metrics import roc_curve, auc

model_lr = LogisticRegression(C=1,class_weight='balanced')
model_lr.fit(x_tfidf_w2v_tr, y_train) # Training Data


y_train_pred = LogisticRegression.predict_proba(model_lr, x_tfidf_w2v_tr)[:,-1] # Predicting Train Data
y_test_pred = LogisticRegression.predict_proba(model_lr, x_tfidf_w2v_te)[:,-1] # Predicting Test Data based on trainin

y_train_pred_con = LogisticRegression.predict(model_lr, x_tfidf_w2v_tr) # Predicting Train Data
y_test_pred_con = LogisticRegression.predict(model_lr, x_tfidf_w2v_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred) # GET TRAIN TRUE POSITIVE and FALSE POSITIVE RATES
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred) # GET TEST TRUE POSITIVE and FALSE POSITIVE RATES

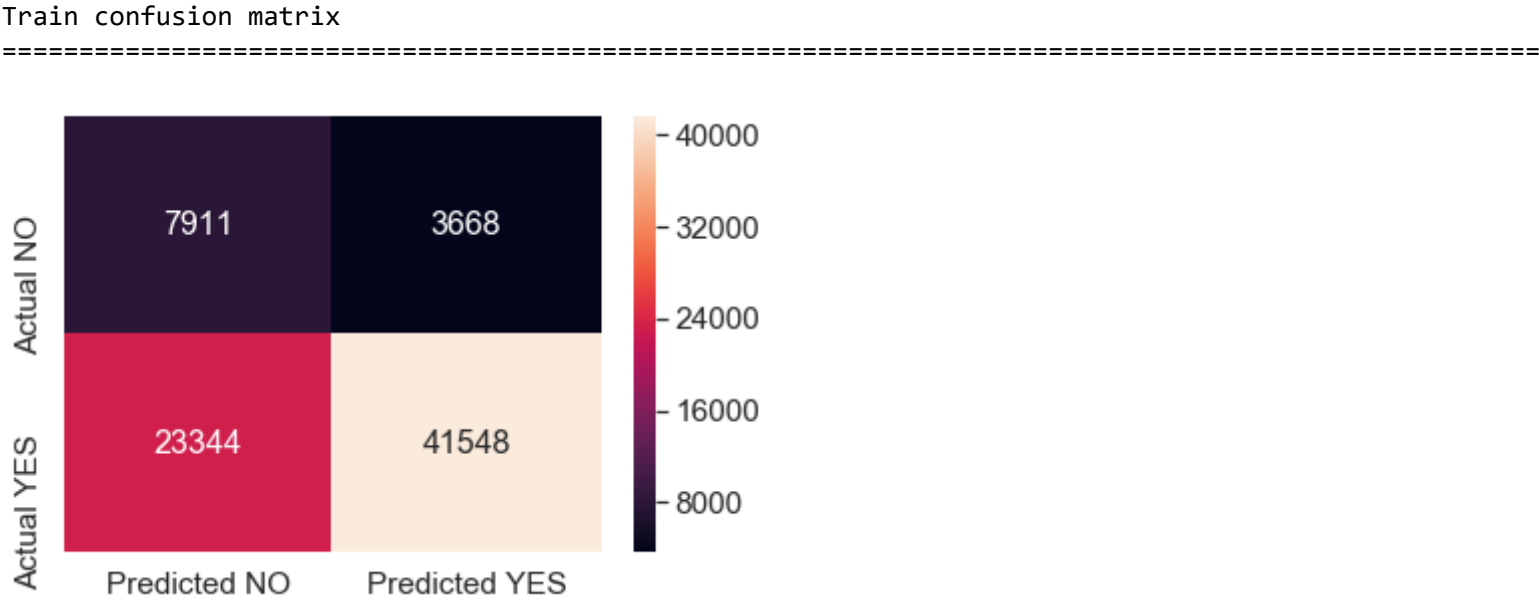
plt.grid()
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS- TFIDF W2v")
plt.show()

auc_tfidf_w2v=auc(test_fpr, test_tpr)
```



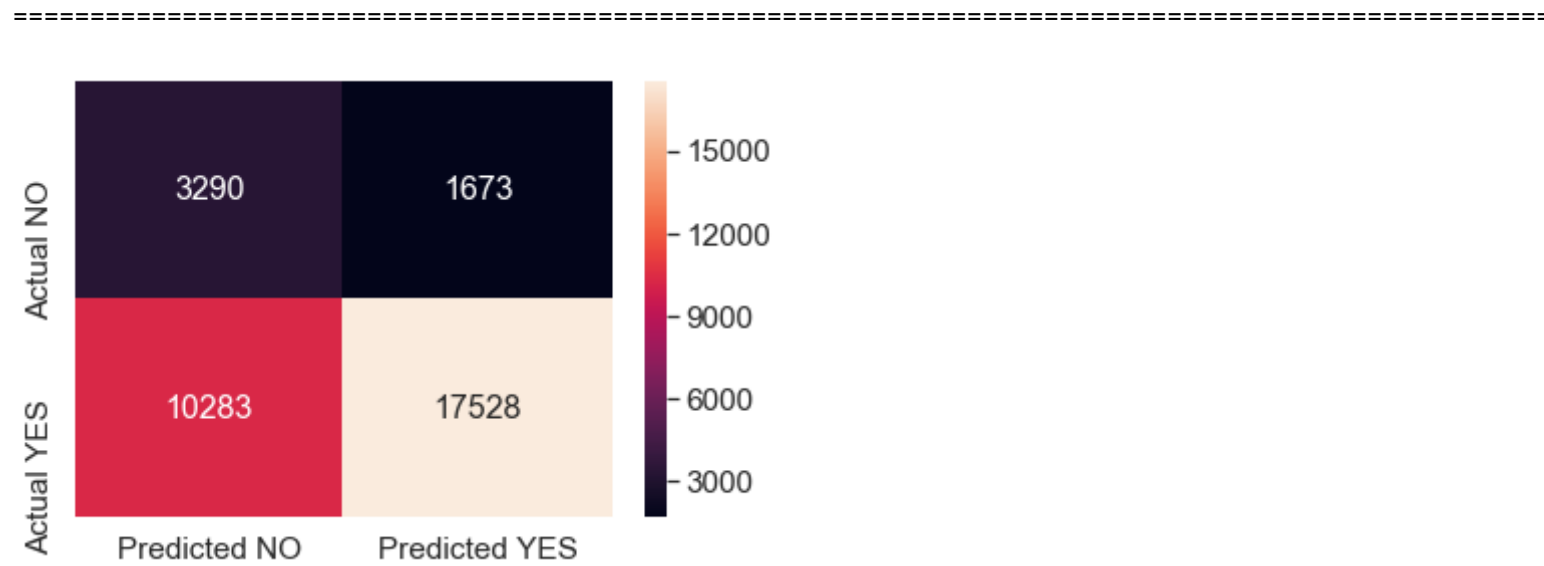
```
In [145]: from sklearn.metrics import confusion_matrix

print("Train confusion matrix")
get_confusion_matrix(y_train,y_train_pred_con)
print('='*100)
```



```
In [146]: print("Test confusion matrix")
print('='*100)
get_confusion_matrix(y_test,y_test_pred_con)
```

Test confusion matrix



2.5 Logistic Regression with added Features Set 5

2.5.1 Get Positive Sentiment score's of each of the essay : numerical data

```
In [147]: from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
analyser = SentimentIntensityAnalyzer()

score_train=[]
for i in x_train['preprocessed_essays'].values:
    score = analyser.polarity_scores(i)
    score_train.append(score.get('pos'))

score_test=[]
for i in x_test['preprocessed_essays'].values:
    score = analyser.polarity_scores(i)
    score_test.append(score.get('pos'))

score_tr=pd.DataFrame(score_train, columns=['pos_sentiment_score'])
score_te=pd.DataFrame(score_test, columns=['pos_sentiment_score'])
```

```
In [148]: # Normalizing numerical features:

#sentiment scores of project essays

from sklearn.preprocessing import Normalizer
normalizer_essays_scores = Normalizer()

x_train_essay_scores_norm = normalizer_essays_scores.fit_transform(score_tr['pos_sentiment_score'].values.reshape(-1,1))
x_test_essay_scores_norm = normalizer_essays_scores.transform(score_te['pos_sentiment_score'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_essay_scores_norm.shape, y_train.shape)
print(x_test_essay_scores_norm.shape, y_test.shape)
print('='*100)
```

After vectorizations

(76471, 1) (76471,)
(32774, 1) (32774,)

2.5.2 Get number of words in the title : numerical data

```
In [149]: title_word_tr=[]
for i in x_train['preprocessed_project_title'].values:
    cnt = len(i.split())
    title_word_tr.append(cnt)

title_word_te=[]
for j in x_test['preprocessed_project_title'].values:
    cnt = len(i.split())
    title_word_te.append(cnt)

title_count_tr=pd.DataFrame(title_word_tr, columns=['count'])
title_count_te=pd.DataFrame(title_word_te, columns=['count'])
```

```
In [150]: # Normalizing numerical features:

#sentiment scores of project essays

from sklearn.preprocessing import Normalizer
normalizer_title_word = Normalizer()

x_train_title_word_norm = normalizer_title_word.fit_transform(title_count_tr['count'].values.reshape(-1,1))
x_test_title_word_norm = normalizer_title_word.transform(title_count_te['count'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_title_word_norm.shape, y_train.shape)
print(x_test_title_word_norm.shape, y_test.shape)
print("=*100)
```

After vectorizations

(76471, 1) (76471,)

(32774, 1) (32774,)

=====

2.5.3 Get number of words in the combine essays : numerical data

```
In [151]: essay_word_tr=[]
for i in x_train['preprocessed_essays'].values:
    cnt = len(i.split())
    essay_word_tr.append(cnt)

essay_word_te=[]
for j in x_test['preprocessed_essays'].values:
    cnt = len(i.split())
    essay_word_te.append(cnt)

essay_count_tr=pd.DataFrame(essay_word_tr, columns=['count'])
essay_count_te=pd.DataFrame(essay_word_te, columns=['count'])
```

```
In [152]: # Normalizing numerical features:

#sentiment scores of project essays

from sklearn.preprocessing import Normalizer
normalizer_essay_word = Normalizer()

x_train_essay_word_norm = normalizer_essay_word.fit_transform(essay_count_tr['count'].values.reshape(-1,1))
x_test_essay_word_norm = normalizer_essay_word.transform(essay_count_te['count'].values.reshape(-1,1))

print("After vectorizations")
print(x_train_essay_word_norm.shape, y_train.shape)
print(x_test_essay_word_norm.shape, y_test.shape)
print("=*100)
```

After vectorizations

(76471, 1) (76471,)

(32774, 1) (32774,)

=====

Combining all features

```
In [153]: #STEP 5

from scipy.sparse import hstack
x_tr = hstack((x_train_state_ohe,x_train_clean_cat_ohe, x_train_clean_subcat_ohe, x_train_grade_ohe, x_train_teacher_ohe, \
               x_train_price_norm, x_train_post_proj_norm, x_train_quantity_norm, \
               x_train_essay_scores_norm,x_train_title_word_norm,x_train_essay_word_norm)).tocsr()

x_te = hstack((x_test_state_ohe,x_test_clean_cat_ohe, x_test_clean_subcat_ohe, x_test_grade_ohe, x_test_teacher_ohe, \
               x_test_price_norm, x_test_post_proj_norm, x_test_quantity_norm, \
               x_test_essay_scores_norm,x_test_title_word_norm,x_test_essay_word_norm)).tocsr()
```

```

In [90]: C=[10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]

tuned_parameters = {'C': C}

tuned_parameters2 = {'C': [math.log(i) for i in C]}

model_lr = RandomizedSearchCV(LogisticRegression(), tuned_parameters, cv=3, scoring='roc_auc')
model_lr.fit(x_tr, y_train)


train_auc= model_lr.cv_results_['mean_train_score']
train_auc_std= model_lr.cv_results_['std_train_score']
cv_auc = model_lr.cv_results_['mean_test_score']
cv_auc_std= model_lr.cv_results_['std_test_score']

plt.plot(tuned_parameters2['C'], train_auc, label='Train AUC')
plt.plot(tuned_parameters2['C'], cv_auc, label='CV AUC')

plt.scatter(tuned_parameters2['C'], train_auc, label='Train AUC points')
plt.scatter(tuned_parameters2['C'], cv_auc, label='CV AUC points')

plt.grid()
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```




```
In [163]: from sklearn.metrics import roc_curve, auc

model_lr = LogisticRegression(C=1,class_weight='balanced')
model_lr.fit(x_tr, y_train) # Training Data


y_train_pred = LogisticRegression.predict_proba(model_lr, x_tr)[:,-1] # Predicting Train Data
y_test_pred = LogisticRegression.predict_proba(model_lr, x_te)[:,-1] # Predicting Test Data based on training model

y_train_pred_con = LogisticRegression.predict(model_lr, x_tr)
y_test_pred_con = LogisticRegression.predict(model_lr, x_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred) # GET TRAIN TRUE POSITIVE and FALSE POSITIVE RATES
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred) # GET TEST TRUE POSITIVE and FALSE POSITIVE RATES

plt.grid()
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

auc=auc(test_fpr, test_tpr)
```




```
In [164]: from sklearn.metrics import confusion_matrix


print("Train confusion matrix")
get_confusion_matrix(y_train,y_train_pred_con)
print('='*100)
```

Train confusion matrix

=====



```
In [165]: print("Test confusion matrix")
get_confusion_matrix(y_test,y_test_pred_con)
print('='*100)
```



3. Conclusion

```
In [166]: #http://zetcode.com/python/prettytable/

from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Model", "Hyper-Parameter", "AUC"]

x.add_row(["BOW", "Brute", 0.01, auc_bow])
x.add_row(["TFIDF", "Brute", 0.1, auc_tfidf])
x.add_row(["AVG-W2V", "Brute", 1, auc_w2v])
x.add_row(["TFIDF-W2V", "Brute", 1, auc_tfidf_w2v])
x.add_row(["NUMERICAL VALUES", "Brute", 1, auc])

print(x)
```

Vectorizer	Model	Hyper-Parameter	AUC
BOW	Brute	0.01	0.6802679477915439
TFIDF	Brute	0.1	0.6867436338603266
AVG-W2V	Brute	1	0.7067704269296581
TFIDF-W2V	Brute	1	0.6980291386130437
NUMERICAL VALUES	Brute	1	0.5734263871588303